

Figure 1: Semantic similarity scores between prompts and their corresponding outputs across different prompt categories.

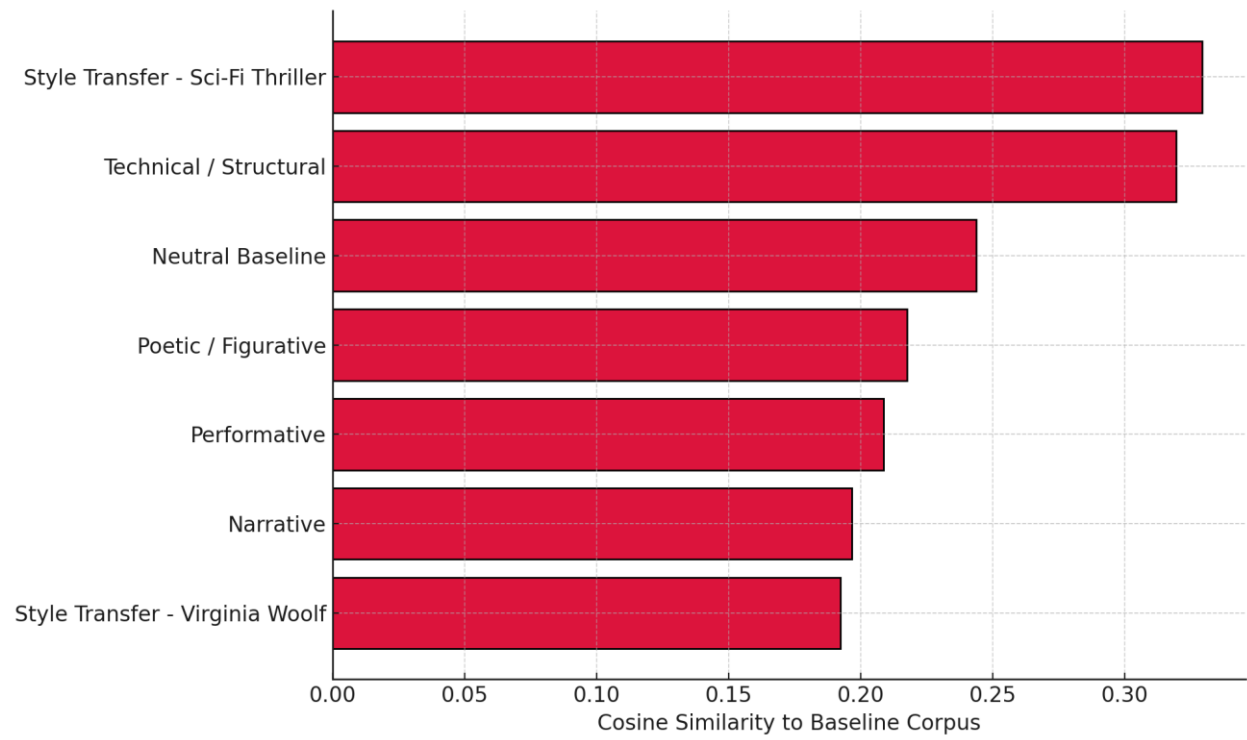
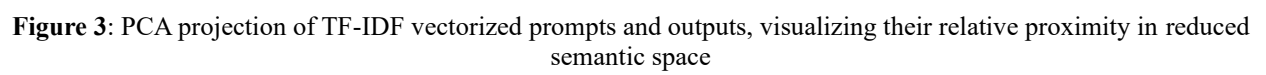


Figure 2: Originality scores of generated outputs compared to a baseline corpus.



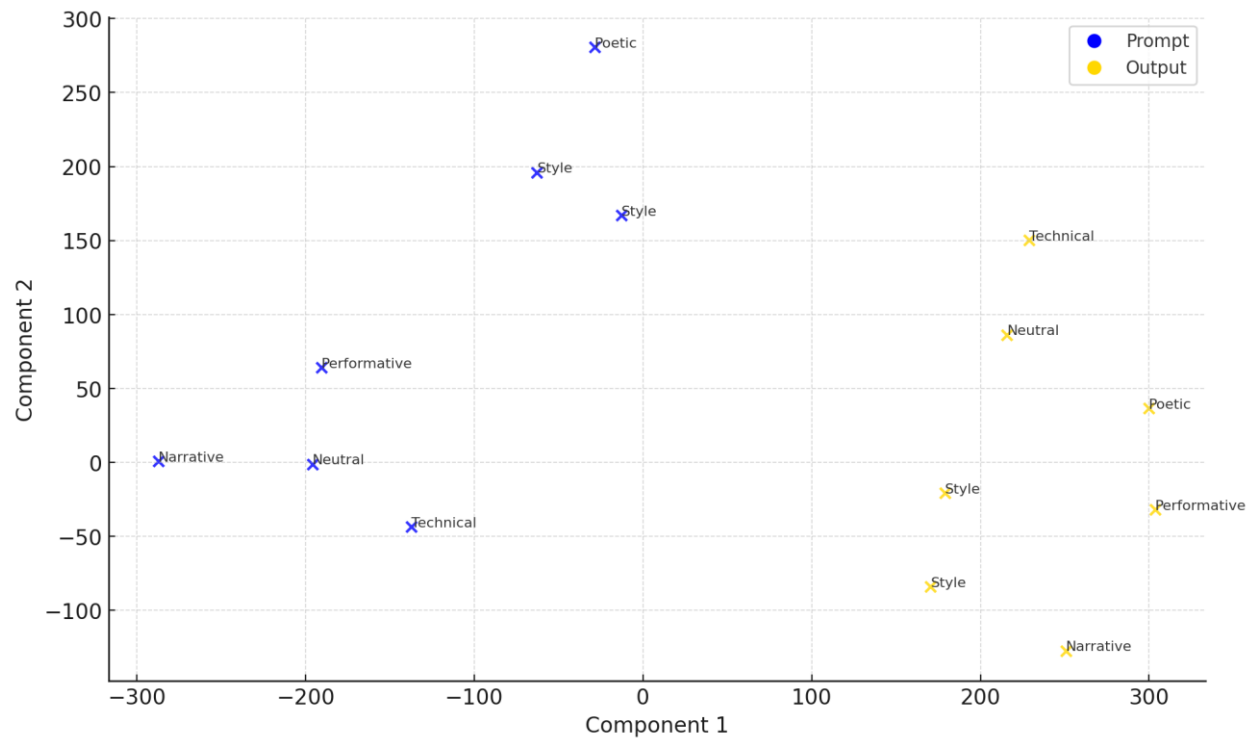


Figure 4: t-SNE projection of TF-IDF embeddings showing semantic drift between prompts and outputs.

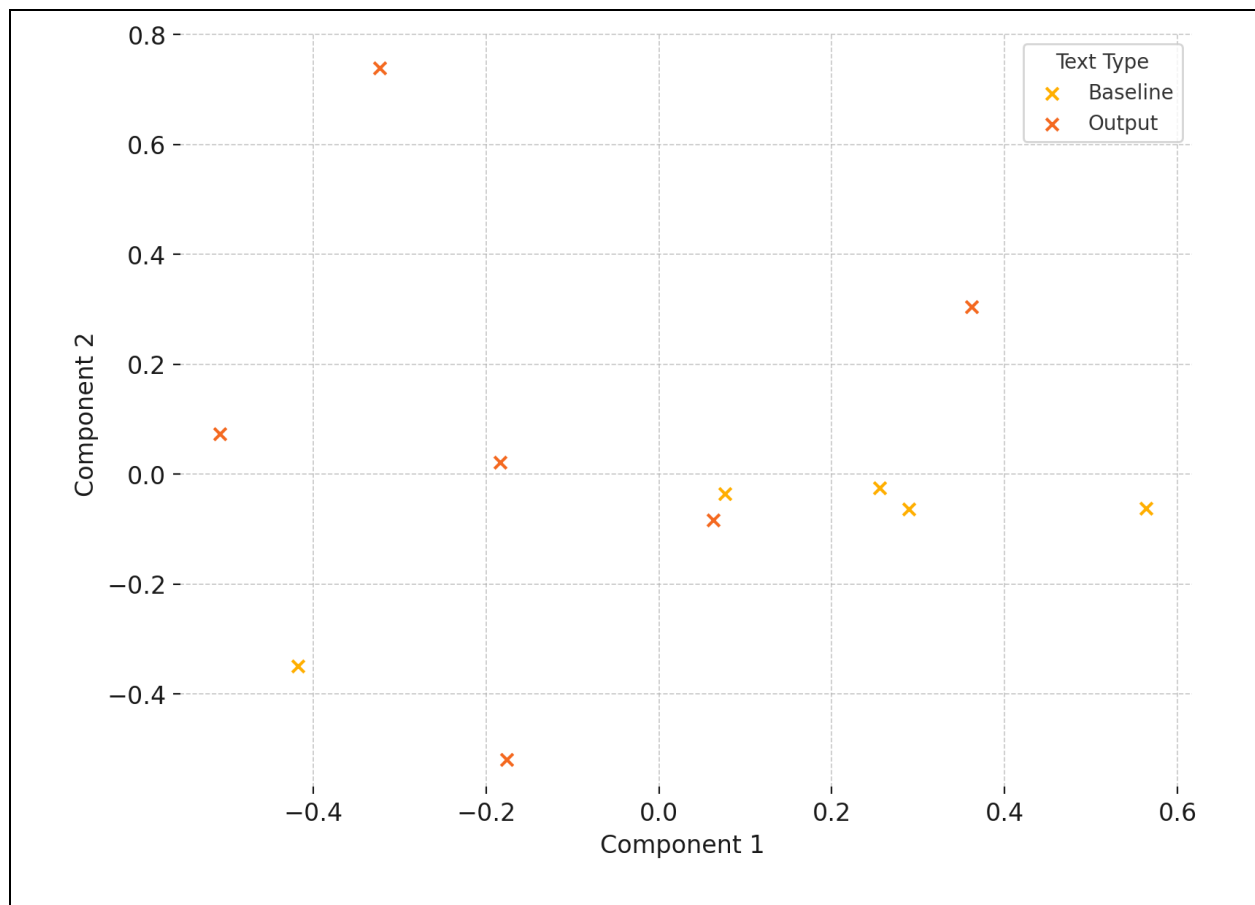


Figure 5: PCA visualization comparing baseline corpus entries and generated outputs, illustrating distribution and divergence in semantic space.

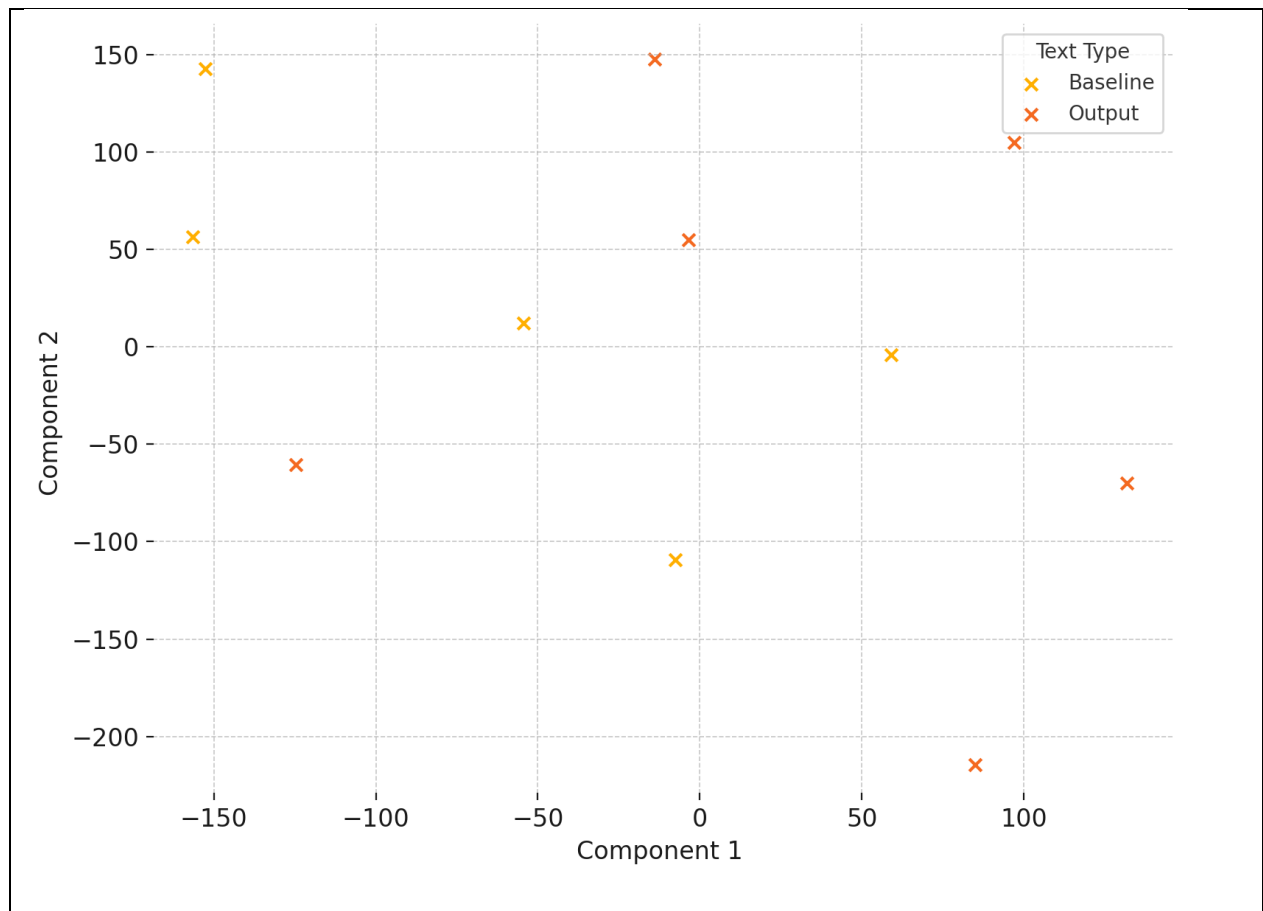


Figure 6: t-SNE visualization of baseline and output texts, highlighting clusters and semantic distance reflective of originality levels.

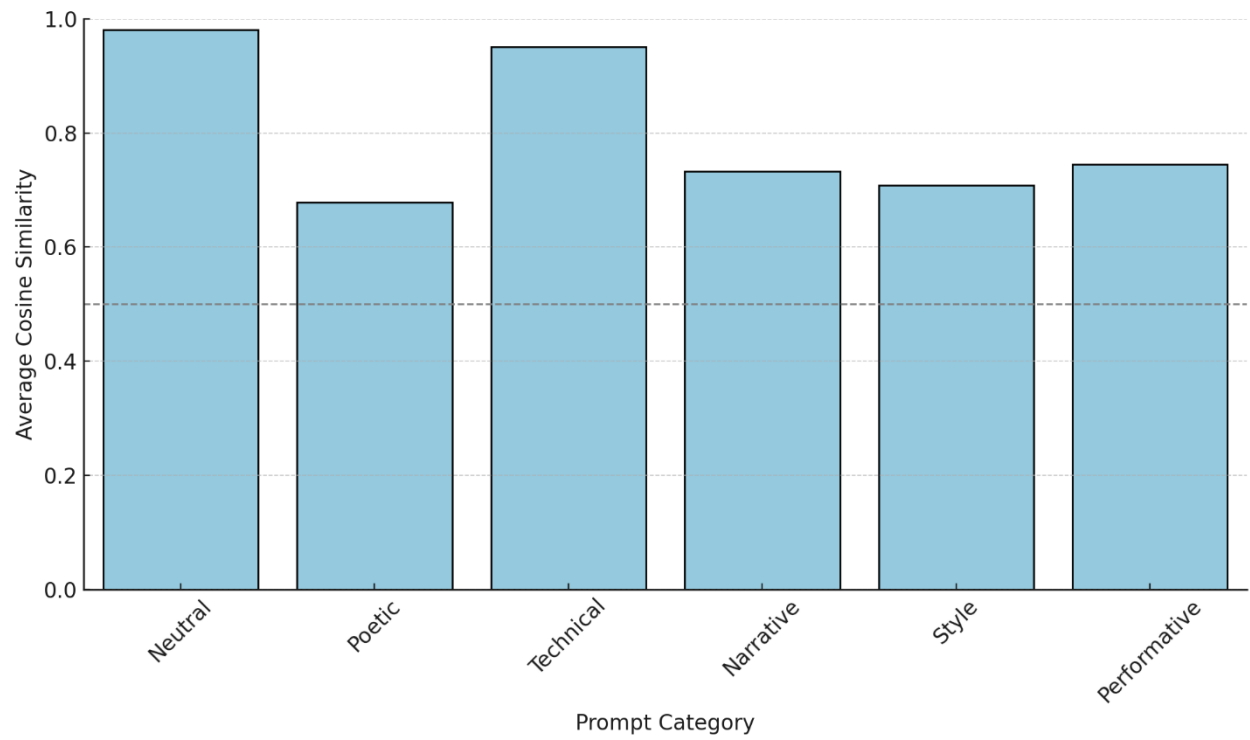


Figure 7: Average cosine similarity between each prompt and its corresponding AI-generated output

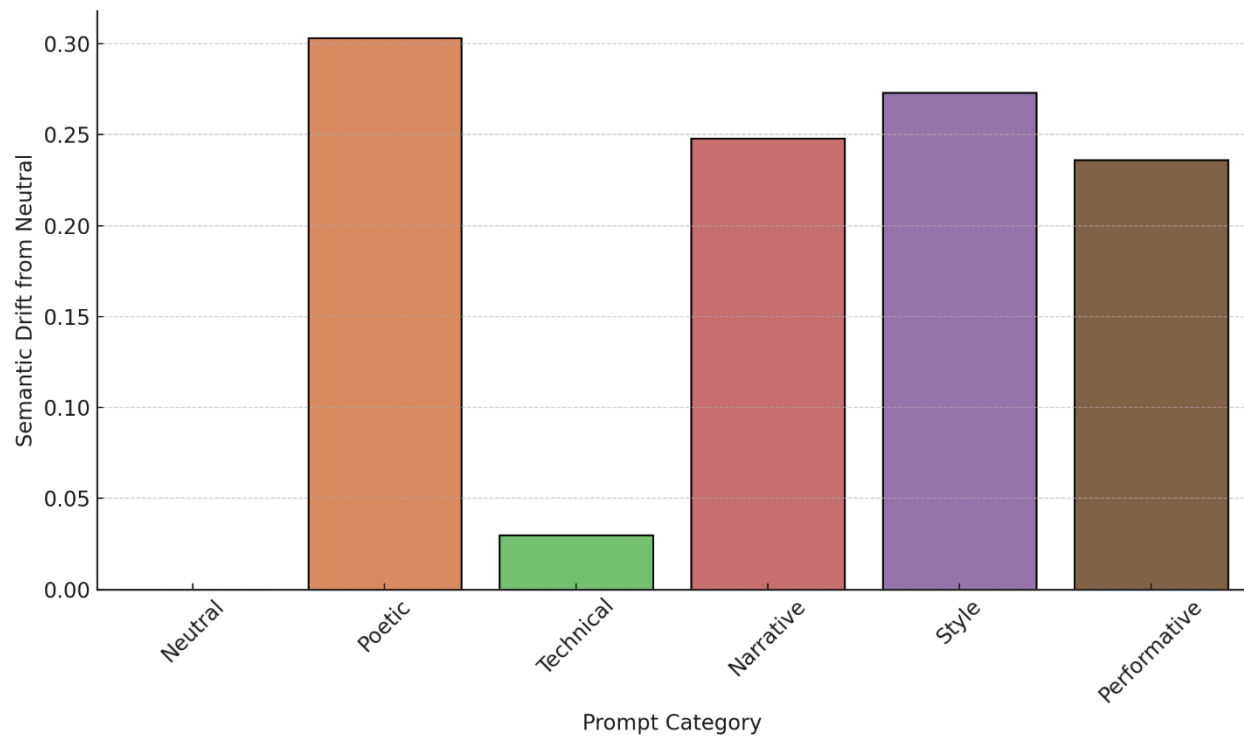


Figure 8: Semantic drift from the Neutral baseline prompt across all prompt categories